

# Efficient RL for Large Language Models

## with Intrinsic Exploration (PREPO)

Yan Sun<sup>1,2</sup>, Jia Guo<sup>2</sup>, Stanley Kok<sup>1</sup>, Zihao Wang<sup>2</sup>, Zujie Wen<sup>2</sup>, Zhiqiang Zhang<sup>2</sup>

<sup>1</sup>National University of Singapore, <sup>2</sup>Ant Group

NeurIPS 2025 Efficient Reasoning Workshop



GitHub




Paper

TL;DR

PREPO reduces Reinforcement Learning (RLVR) training costs


with "intrinsic" metrics—Prompt Perplexity & Rollout Entropy—to filter data, guiding exploration.

### I. The Problem



- Costly:** Standard RLVR generates thousands of rollouts.
- Inefficient:** Many samples are too easy or too hard (zero advantage).
- Goal:** Data-efficient RLVR training using data intrinsic properties.


### 2. Preliminary Analysis



### 3. Method: Online Prompt Selection

#### Strategy: Prompt Perplexity

Use perplexity as a proxy to select from a candidate batch  $\mathcal{B}$  to the actual batch  $\mathcal{I}_p$  at every training step. Train on Low PPL to High PPL prompts.




### 4. Method: Rollout Weighting

#### Strategy: Relative Entropy

Prioritize diverse reasoning paths. Weight rollouts by their average token-level entropy. ( $V$ : vocabulary size)

$$\bar{H}_i = -\frac{1}{|o_i|} \sum_{t=1}^{|o_i|} \sum_{v \in V} p(v|x_t) \log p(v|x_t)$$

mini-batch  
( $B$  rollouts)



### 5. Results

Tested on Qwen & Llama (MATH500, AIME, Olympiad)

Model	Method	Avg Acc.	Rollouts
Qwen2.5-Math-7B	Random	39.45%	905K
	PREPO	39.59%	540K (1.7x)
Qwen3-4B	Random	71.33%	553K
	PREPO	75.99%	348K (1.6x)

